

IEEE Signal Processing Society Seasonal School

28.03.2022 – 01.04.2022

Networked Federated Learning

Theory, Algorithms and Applications

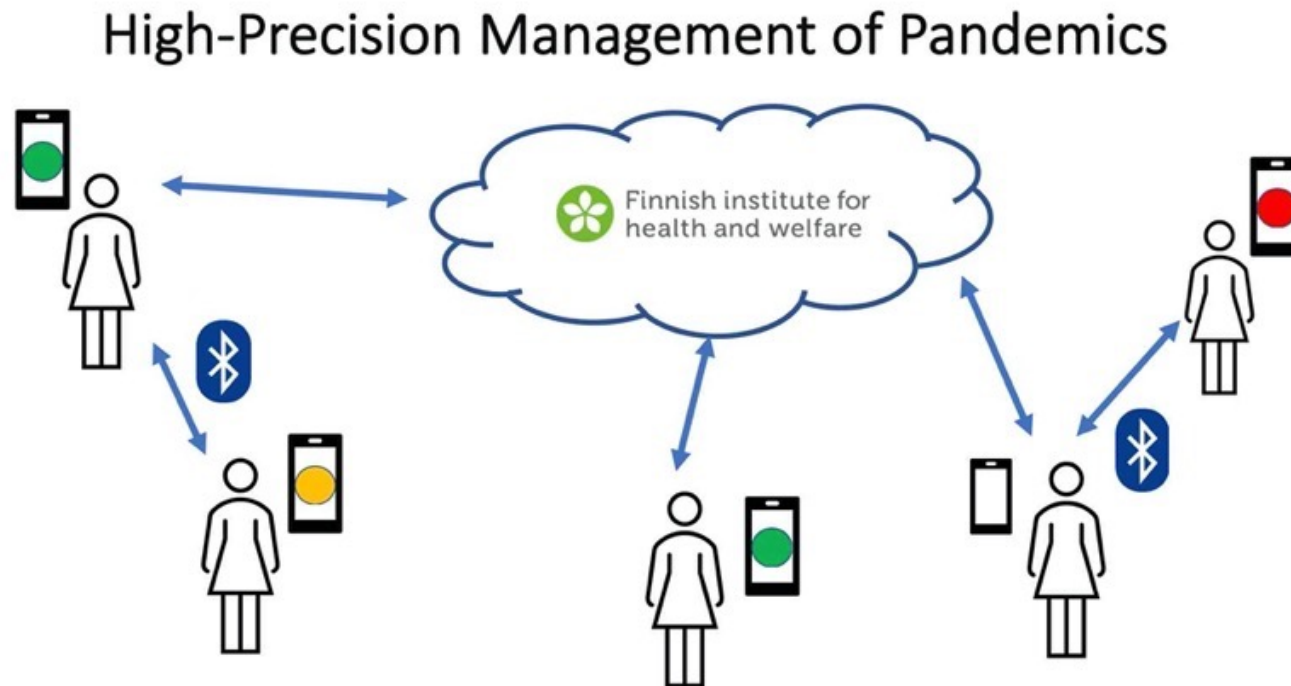
Alex Jung (Aalto University)



About Me.

- MSc (2008) and Ph.D. (2012) in EE, TU Vienna
- since 2015 Ass. Prof. for Machine Learning at Aalto/CS
- leading group “Machine Learning for Big Data”
- 2020- Associate Editor for IEEE Signal Processing Letters
- 2019-2022: Chair of IEEE Finland Jt. Chapter SP/CAS

RA1: Networked Federated Learning.




Y. Sarcheshmehpour, M Leinonen and AJ, "Federated Learning From Big Data Over Networks", IEEE ICASSP, 2021.

AJ, "Networked Exponential Families for Big Data Over Networks," in IEEE Access, 2020, doi: 10.1109/ACCESS.2020.3033817.

AJ, N. Tran, "Localized Linear Regression in Networked Data," in IEEE SPL, 2019, doi: 10.1109/LSP.2019.2918933.

RA2: Explainable Machine Learning.



explanation can be:


- relevant example of training set
- subset of features
- counterfactuals
- a free text explanation
- court sentence

L. Zhang, G. Karakasidis, A. Odnoblyudova, L. Dogruel and AJ., “Explainable Empirical Risk Minimization”, arXiv, 2022. [weblink](#)
AJ and P. H. J. Nardelli, “An Information-Theoretic Approach to Personalized Explainable Machine Learning,” in IEEE SPL, 2020, doi: 10.1109/LSP.2020.2993176.


School Format.



Lectures
(via zoom, recorded)



Exercises
(independent work)



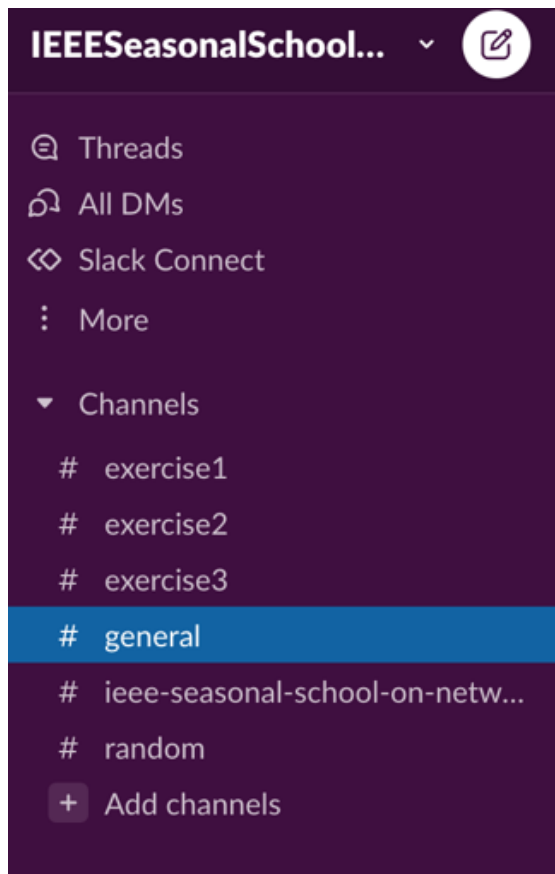
Discussion Forum
[click here to join](#)

all times EET/local Helsinki time)

hall: <https://aalto.zoom.us/j/64034033878?pwd=ayt4aDNyaEVSSGt1eHNNZFVFR2RoZz09>

- Tue, 11:00-13:00, [Basics of spectral graph theory](#), Prof. Avratchenkov
- We, 10:00-11:00, [Parallel/distributed methods for state-space models](#), Prof. Särkkä
- We, 12:30-13:30, [Compact and Efficient Neural Networks, Steps Towards Communication Efficient Federated Learning](#), Dr. H. Tavakoli
- Thu., 11:00-12:00, [Machine Learning applications in meteorological forecasting](#), Dr. Schicker
- Thu, 13:00-14:00, [Towards Communication-Efficient and Personalized Federated Learning](#), Dr. Samek
- Thu, 14:30-15:30, [Communication-Computation Efficient Distributed Machine Learning](#), Prof. Fischione
- Fr, 11:00-12:00, [Tackling the problem of “bad” explanations with the Human-in-the-Loop principle](#), Dipl.-Ing. Saranti

Discussion Forum („Slack“)



can be used via web-browser or with an installed app;

link to join is on school site

<https://ieeespcasfinland.github.io/>

post your questions there in suitable channel, e.g., “exercise1”

don't hesitate to answer other's questions

Teacher Support for Exercises



Dr. Yu Tian




Dr. Shamsiat (Shamsi) Abdurakhmanova

Exercises.

- **three exercises**: Exercise 1, 2 and 3
- each exercise consists of a **Python notebook**
- notebook contains **starter code** and explanations
- **tasks at the end** of notebooks

Running Toy Application



morning temperature = - 10 (minimum daytime temp.)
maximum daytime temperature ?


Exercise 1.



- basic components of ML: data, model and loss
- principle of empirical risk minimization (ERM)
- solve ERM using gradient descent
- read in datapoints from a csv file



Three Components of ML



Data.

```
import pandas as pd # the pandas package provides tools for data manipulation
df = pd.read_csv('FMIDData.csv') # read in data from a csv file
df.head(5) # show first 5 datapoints
```

	date	min_temp	max_temp	latitude	longitude	station
0	2021-04-11	x ¹⁴ (2.4)	y ⁽⁰⁾ 6.1	60.30373	25.54916	Porvoo Kilpilahti satama
1	2021-04-12	4.9	9.9	60.30373	25.54916	Porvoo Kilpilahti satama
2	2021-04-13	2.8	6.5	60.30373	25.54916	Porvoo Kilpilahti satama
3	2021-04-14	0.5	6.5	60.30373	25.54916	Porvoo Kilpilahti satama
4	2021-04-15	0.7	9.4	60.30373	25.54916	Porvoo Kilpilahti satama

Feature Matrix and Label Vector.

	date	min_temp	max_temp
0	2021-04-11	2.4	6.1
1	2021-04-12	4.9	9.9
2	2021-04-13	2.8	6.5
3	2021-04-14	0.5	6.5
4	2021-04-15	0.7	9.4

$\Rightarrow y \in \mathbb{R}^5$

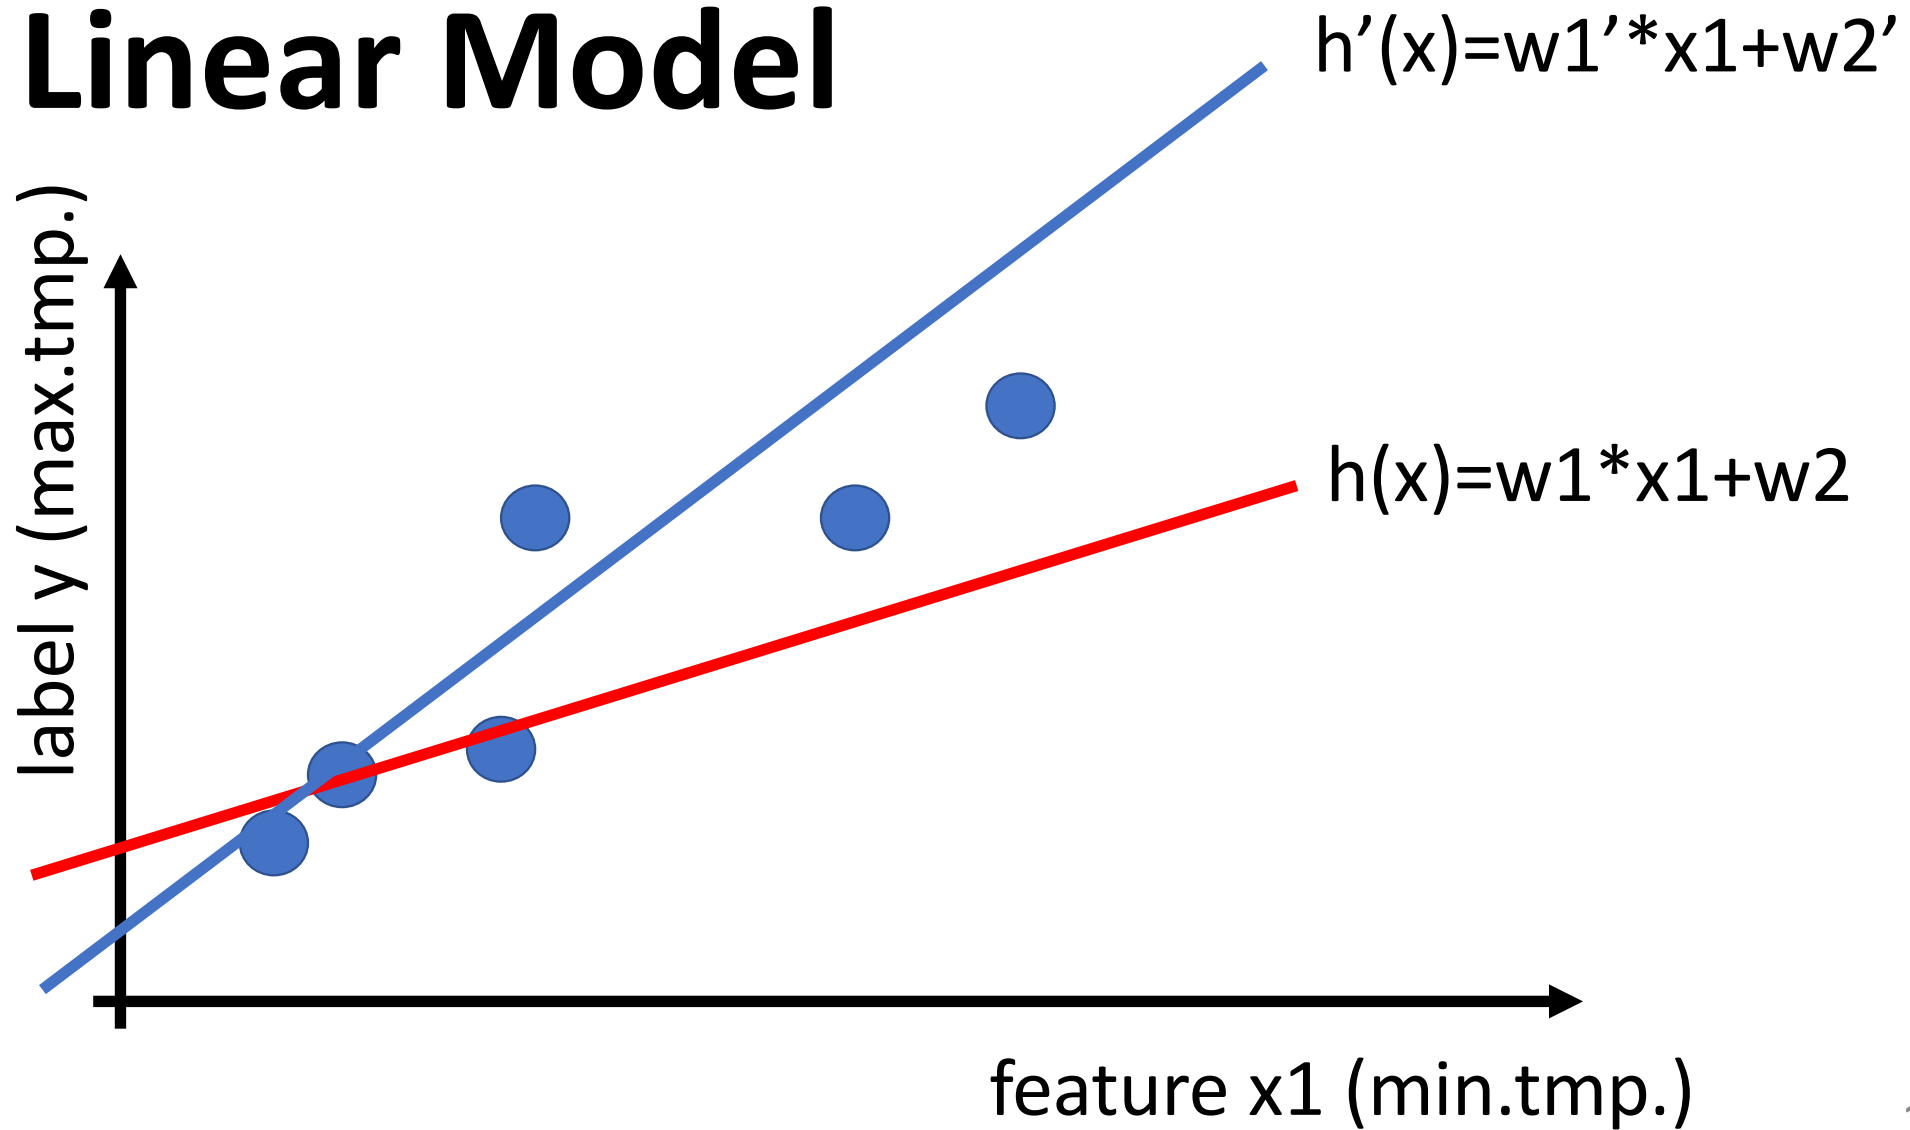
$X = \begin{pmatrix} 2.4 & 1 \\ 4.9 & 1 \\ \vdots & \vdots \\ 0.7 & 1 \end{pmatrix}$

$x^{(1)} = (x_1^{(1)} \ x_2^{(1)})$

$x^{(2)}$

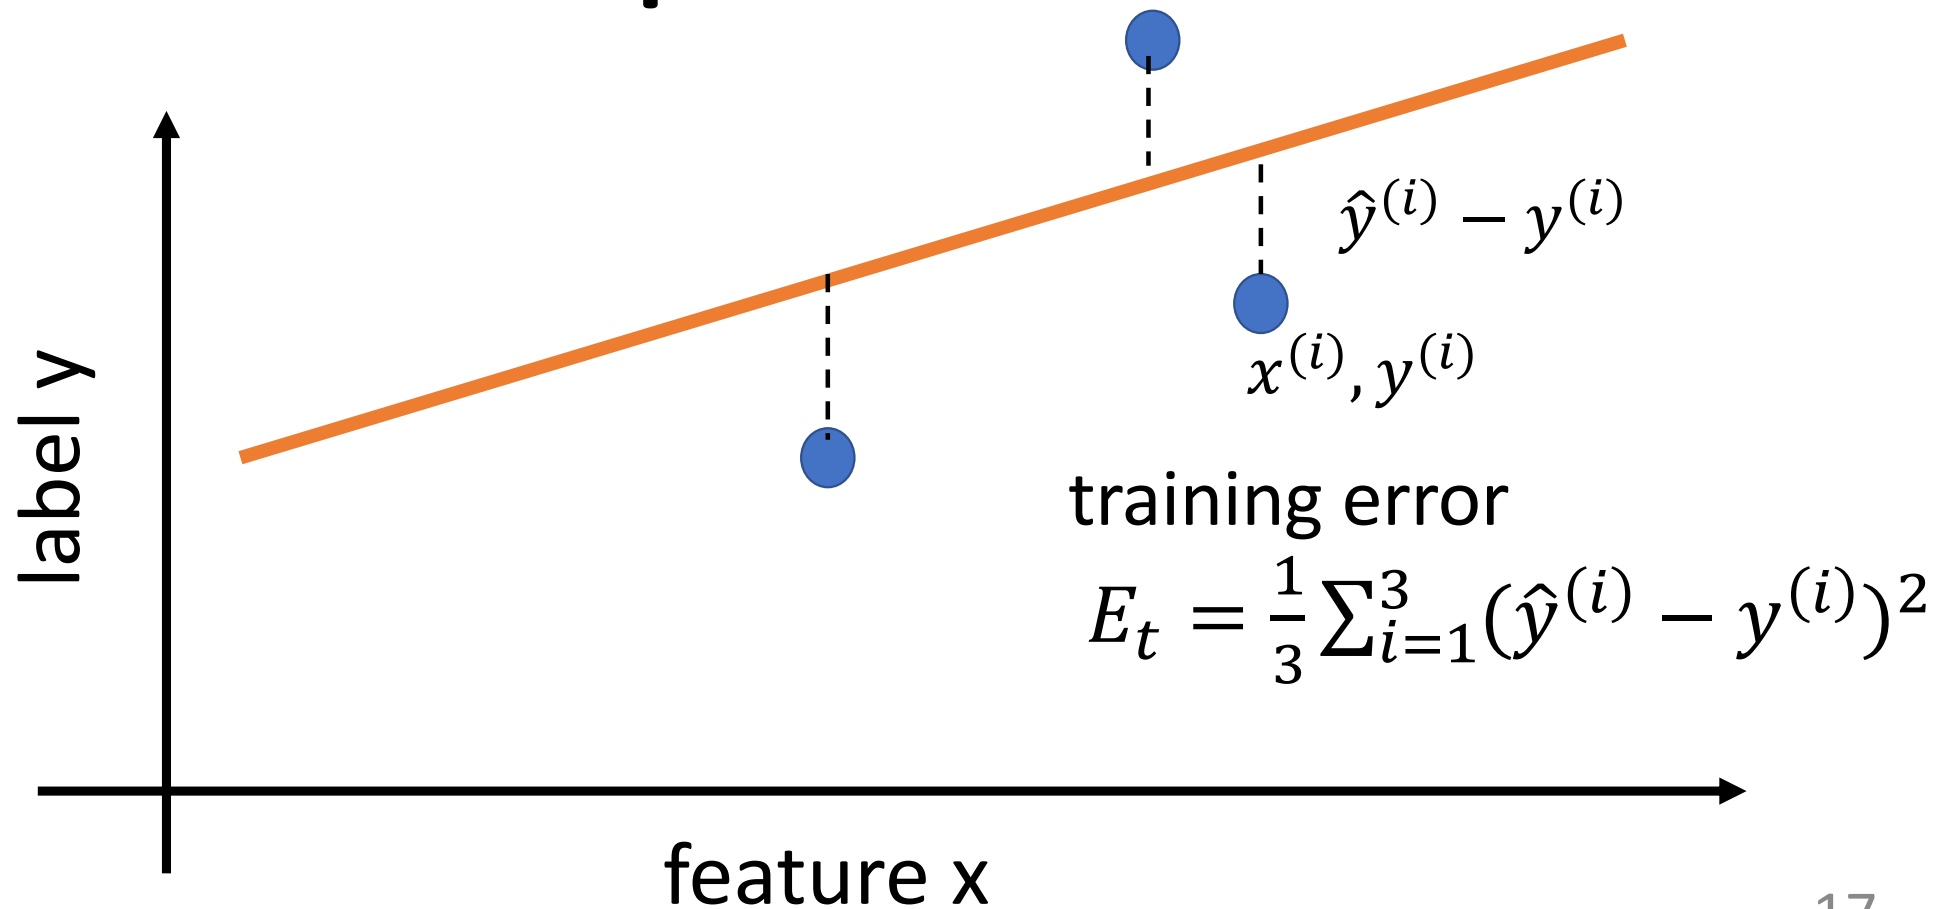
```
X = np.hstack((df["min_temp"].to_numpy().reshape(m,1), np.ones((m,1))))  
y = df["max_temp"].to_numpy().reshape(m,)
```

Linear Model



which hypothesis $h(x)$, or
model parameter w , is best?

Mean Squared Error



Empirical Risk Minimization (informal)

*learn hypothesis (model parameters) such that
the resulting average loss (“empirical risk”)
on a training set is minimal*

Empirical Risk Minimization

datapoints $(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})$,

$$\min_{\mathbf{w} \in \mathbb{R}^n} (1/m) \sum_{i=1}^m \left(\underbrace{\mathbf{w}^T \mathbf{x}^{(i)}}_{h(\mathbf{x}^{(i)})} - y^{(i)} \right)^2.$$

Matrix/Vector Notation

$$\min_{\mathbf{w}} f(\mathbf{w}) , \text{ with } f(\mathbf{w}) := (1/m) \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$

feature matrix $\mathbf{X} = \left(\left(\mathbf{x}^{(1)} \right)^T, \dots, \left(\mathbf{x}^{(m)} \right)^T \right)^T$

label vector $\mathbf{y} = \left(y^{(1)}, \dots, y^{(m)} \right)^T.$

Gradient Descent

$$\min_{\mathbf{w}} f(\mathbf{w}) , \text{ with } f(\mathbf{w}) := (1/m) \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$


start with some initial guess $\mathbf{w}^{(0)}$ and iteratively improve by GD steps

$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \alpha \nabla f(\mathbf{w}^{(r)})$$


with step-size or “learning rate” α

Gradient Descent Step


$$\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \alpha \nabla f(\mathbf{w}^{(r)})$$



Effect of Learning-Rate.



α too small: GD steps make only very little progress!



α too large: GD steps depart from optimum!

A Sufficient Condition.

assume objective function is “ β -smooth”

$$\|\nabla f(\mathbf{w}) - \nabla f(\mathbf{w}')\| \leq \beta \|\mathbf{w} - \mathbf{w}'\|.$$

GD converges when using learning rate

$$\alpha = 1/\beta$$

proof: <https://www.stat.cmu.edu/~ryantibs/convexopt-F13/scribes/lec6.pdf>

Computing Gradient.

to implement GD step $\mathbf{w}^{(r+1)} = \mathbf{w}^{(r)} - \alpha \nabla f(\mathbf{w}^{(r)})$

we need to compute gradient $\nabla f(\mathbf{w}^{(r)})$

for $f(\mathbf{w}) := (1/m) \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$, we obtain

$$\nabla f(\mathbf{w}) = -(2/m) \mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

see A.4 of S. Boyd and L. Vandenberghe, "Convex Optimization", CUP, 2004

GD for Linear Regression.


$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \alpha(2/m)\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}^{(k)})$$

```
# compute the gradient of f(w) at the current weight vector (ok
gradient = -(2/m) * X.T.dot(y - X.dot(current_weights))
# update the current weight vector via the GD step
current_weights = current_weights - (learning_rate * gradient)
```


When to stop ?

- nr. of steps required to ensure sub-optimality ("theory")
- check for sufficiently large decrease $f(\mathbf{w}^{(k+1)}) - f(\mathbf{w}^{(k)})$
- try $h(x) = \mathbf{w}^{(k)}x$ on validation set ("early stopping")

GD in Action.



Train and Validate.



Exercise 1 – Task 1.1

Generate synthetic dataset of m data points. Each datapoint characterized by feature vector $\mathbf{x} = (x_1, x_2)^T$ with $x_1 \sim \mathcal{N}(0,1)$ and $x_2 = 1$. and numeric label $y = \bar{w}_1 x_1 + \bar{w}_2 x_2 + \sigma \varepsilon$. with some given “true weights” \bar{w}_1, \bar{w}_2 a given noise strength σ and Gaussian noise $\varepsilon \sim \mathcal{N}(0,1)$

Study the deviation of the learnt weight vector after a given number N of GD steps from the true weight vector (\bar{w}_1, \bar{w}_2) as a function of N, σ and m .

Exercise 1 – Task 1.2

Learn the weights of a linear predictor for the maximum daytime temperature (label) from the minimum daytime temperature (feature) at some place in Finland. To achieve this goal, use the datapoints in

<https://raw.githubusercontent.com/ieeespcasfinland/ieeespcasfinland.github.io/main/FMIDData.csv>

Questions ?

5 Min Break

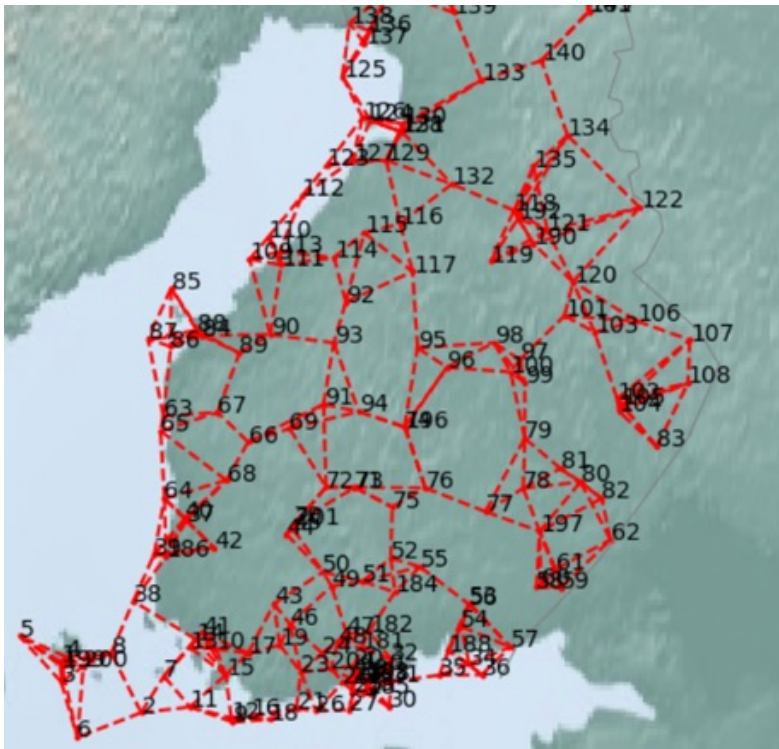


Exercise 2.



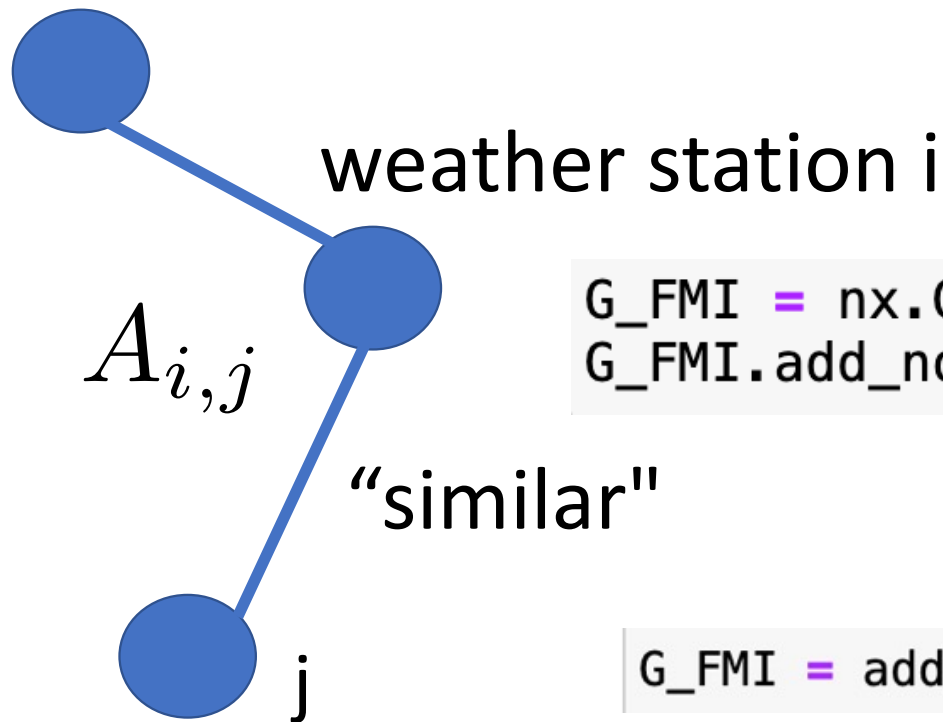
- store networked data and models using [networkx.Graph](#) object
- add local datasets and model parameters as node attributes
- combine GD with a network averaging algorithm to collaboratively linear hypothesis from a network of local datasets

Network of Weather Data



2021-04-24, 1.2, 5.8, 60.30373, 25.54916, Porvoo Kilpilanti satama
 2021-04-25, 1.0, 6.5, 60.30373, 25.54916, Porvoo Kilpilahti satama
 2021-04-26, 1.4, 5.1, 60.30373, 25.54916, Porvoo Kilpilahti satama
 2021-04-27, 1.5, 5.3, 60.30373, 25.54916, Porvoo Kilpilahti satama
 2021-04-28, 0.2, 6.5, 60.30373, 25.54916, Porvoo Kilpilahti satama
 2021-04-29, -1.2, 7.7, 60.30373, 25.54916, Porvoo Kilpilahti satama
 2021-04-30, -1.0, 9.0, 60.30373, 25.54916, Porvoo Kilpilahti satama
 2021-04-11, -1.7, 7.7, 60.12735, 19.90038, Jomala Maarianhamina lentoasema
 2021-04-12, 3.9, 7.5, 60.12735, 19.90038, Jomala Maarianhamina lentoasema
 2021-04-13, -0.2, 5.2, 60.12735, 19.90038, Jomala Maarianhamina lentoasema
 2021-04-14, -3.3, 6.0, 60.12735, 19.90038, Jomala Maarianhamina lentoasema
 2021-04-15, -4.4, 9.0, 60.12735, 19.90038, Jomala Maarianhamina lentoasema
 2021-04-16, -0.2, 11.5, 60.12735, 19.90038, Jomala Maarianhamina lentoasema
 2021-04-17, 0.3, 14.5, 60.12735, 19.90038, Jomala Maarianhamina lentoasema

The Empirical Graph.




```
G_FMI = nx.Graph()  
G_FMI.add_nodes_from(range(0, num_stations))
```

```
G_FMI = add_edges(G_FMI, total_neigh=4)
```

edge weights $A_{i,j}$ quantify
"statistical similarities"

Attaching Local Datasets to Nodes

weather station i


$$\mathbf{X}^{(i)} = \left(\mathbf{x}^{(i,1)}, \dots, \mathbf{x}^{(i,m_i)} \right)^T, \text{ and } \mathbf{y}^{(i)} = \left(y^{(i,1)}, \dots, y^{(i,m_i)} \right)^T.$$

```
G_FMI.nodes[i]['X'] = df.min_temp.to_numpy().reshape(-1,1)
G_FMI.nodes[i]['y'] = df.max_temp.to_numpy() # label vector
```

Attaching Linear Model to Nodes

weather station i

$$\mathbf{X}^{(i)} = (\mathbf{x}^{(i,1)}, \dots, \mathbf{x}^{(i,m_i)})^T, \text{ and } \mathbf{y}^{(i)} = (y^{(i,1)}, \dots, y^{(i,m_i)})^T.$$

$$h^{(i)}(\mathbf{x}) = (\mathbf{w}^{(i)})^T \mathbf{x}$$

```
G.nodes[iter_node]["w"] = np.zeros(X.shape[1])
```

```
G_FMI.nodes[i]['X'] = df.min_temp.to_numpy().reshape(-1,1)  
G_FMI.nodes[i]['y'] = df.max_temp.to_numpy() # label vector
```

Learning a Global Model

$$\mathbf{w}^{(0)} = \mathbf{w}^{(1)} = \dots = \mathbf{w}^{(m-1)} = \mathbf{w}$$

learn weight vector \mathbf{w} via GD applied to linear regression with pooling all local datasets


$$\mathbf{X} = \left(\left(\mathbf{X}^{(0)} \right)^T, \dots, \left(\mathbf{X}^{(m-1)} \right)^T \right)^T, \mathbf{y} = \left(\left(\mathbf{y}^{(0)} \right)^T, \dots, \left(\mathbf{y}^{(m-1)} \right)^T \right)^T$$

Centralized Federated Learning

GD step for linear regression on pooled local datasets:

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha(1/m) \sum_{i=0}^{m-1} g^{(i)}(\mathbf{w}^{(k)})$$

local gradient $g^{(i)}(\mathbf{w}) := -2(\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)}\mathbf{w})$



Distributed Federated Learning

GD step for linear regression on pooled local datasets:


$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha(1/m) \sum_{i=0}^{m-1} g^{(i)}(\mathbf{w}^{(k)})$$

local gradient $g^{(i)}(\mathbf{w}) := -2(\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)}\mathbf{w})$

use network averaging to (approximately) compute

$$(1/m) \sum_{i=0}^{m-1} g^{(i)}(\mathbf{w}^{(k)})$$

Network Averaging



$$\mathbf{g}^{(i)}(r+1) = W_{i,i}\mathbf{g}^{(i)}(r) + W_{i,j}\mathbf{g}^{(j)}(r) + W_{i,k}\mathbf{g}^{(k)}(r)$$

for suitable choice of weights $W_{i,j}$,

$$\lim_{r \rightarrow \infty} \mathbf{g}^{(i)}(r+1) = (1/m) \sum_{i'=0}^{m-1} \mathbf{g}^{(i')}(0)$$

Algorithm 1

1. init weights to zero at all nodes

```
G.nodes[iter_node]["w"] = np.zeros(X.shape[1])
```

2. repeat for N_GD times:

2.1 compute local gradients at all nodes

```
# compute local gradient for current weight vector at node iter_node  
G.nodes[iter_node]["g"] = -2*X.T.dot(y - X.dot(G.nodes[iter_node]["w"]))
```

2.2 N_AV iterations of network averaging

2.3 do a GD step at each node

```
for iter_node in G.nodes(data=False):  
    G.nodes[iter_node]["w"] = G.nodes[iter_node]["w"] - (learning_rate * G.nodes[iter_node]["g"] )
```

Exercise 2 – Task 2.1

Try out Algorithm 1 for a toy networked dataset consisting of two clusters/blocks

study effect of having an edge between clusters, varying number of GD steps, varying number of network averaging iterations

Exercise 2 – Task 2.2

Try out Algorithm 1 for a networked data obtained from Finnish meteorological institute (Demo code shows to load this data)

Questions ?

5 Min Break



Exercise 3.

- couple local models using total variation (TV)
- gradient descent for TV regularized local linear regression

Networked Data and Models.

weather station i


$$\mathbf{X}^{(i)} = (\mathbf{x}^{(i,1)}, \dots, \mathbf{x}^{(i,m_i)})^T, \text{ and } \mathbf{y}^{(i)} = (y^{(i,1)}, \dots, y^{(i,m_i)})^T.$$

$$h^{(i)}(\mathbf{x}) = (\mathbf{w}^{(i)})^T \mathbf{x}$$

```
G.nodes[iter_node]["w"] = np.zeros(X.shape[1])
```

```
G_FMI.nodes[i]['X'] = df.min_temp.to_numpy().reshape(-1,1)  
G_FMI.nodes[i]['y'] = df.max_temp.to_numpy() # label vector
```

Learn Personalized Models




Local Linear Regression

$$\min_{\mathbf{w}^{(i)} \in \mathbb{R}^n} (1/m_i) \left\| \mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)} \right\|_2^2.$$


we could learn local weight vector by solving
local linear regression for each node

not a good idea if local dataset is too small

Clustering Assumption



Total Variation (TV)



requiring small TV

$$\sum_{\{i,j\}} A_{i,j} \|\mathbf{w}^{(i)} - \mathbf{w}^{(j)}\|^2$$

enforces weights to be nearly constant over clusters

TV-Regularized Linear Regression

$$\min_{\mathbf{w}^{(0)}, \dots, \mathbf{w}^{(i)}} \sum_i (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2 + \lambda_{\text{TV}} \sum_{\{i,j\}} A_{i,j} \|\mathbf{w}^{(i)} - \mathbf{w}^{(j)}\|_2^2.$$

increasing λ_{TV}

“clusteredness”

local training errors


Gradient Descent

$$\min_{\mathbf{w}^{(0)}, \dots, \mathbf{w}^{(i)}} \sum_i (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2 + \lambda_{\text{TV}} \sum_{\{i,j\}} A_{i,j} \|\mathbf{w}^{(i)} - \mathbf{w}^{(j)}\|_2^2.$$

objective function is smooth and convex;
-> can be solved iteratively using GD steps

$$\mathbf{w}^{(i,k+1)} = \mathbf{w}^{(i,k)} - \alpha \left(-(2/m_i) (\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i,k)}) + 2\lambda_{\text{TV}} \sum_{i' \neq i} A_{i,i'} (\mathbf{w}^{(i,k)} - \mathbf{w}^{(i',k)}) \right)$$

GD Step as Message Passing



$$\mathbf{w}^{(i,k+1)} = \mathbf{w}^{(i,k)} - \alpha \left(-(2/m_i) (\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i,k)}) + 2\lambda_{\text{TV}} \sum_{i' \neq i} A_{i,i'} (\mathbf{w}^{(i,k)} - \mathbf{w}^{(i',k)}) \right)$$

Algorithm 2

1. init weights to zero all nodes

```
G.nodes[iter_node]["w"] = np.zeros(X.shape[1])
```

2. repeat for N_GD times:

update local weights at all nodes as

$$\mathbf{w}^{(i,k+1)} = \mathbf{w}^{(i,k)} - \alpha \left(-(2/m_i) (\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i,k)}) + 2\lambda_{\text{TV}} \sum_{i' \neq i} A_{i,i'} (\mathbf{w}^{(i,k)} - \mathbf{w}^{(i',k)}) \right)$$

Exercise 3 – Task 3.1

Try out Algorithm 2 for a toy networked dataset consisting of two clusters/blocks

study the effect of varying edge weight between clusters, varying number of GD steps, varying regularization parameter λ_{TV}

Exercise 3 – Task 3.2

Try out Algorithm 2 for a networked data obtained from Finnish meteorological institute (Demo code shows to load this data)

Thank You!

Enjoy the School Week!